

SUDISH KUMAR

Hyderabad, India | sudishnit@gmail.com | +91 9870176701 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

MACHINE LEARNING ENGINEER | DEEP LEARNING (PYTORCH) | GENERATIVE AI | AWS SAGEMAKER

AWS-certified AI Practitioner & Generative AI Developer with 10+ years of professional experience, including 3+ years building AI/ML systems at Amazon. Specialized in production agentic AI, RAG pipelines, LLM evaluation, and full-lifecycle MLOps. Maintained 99%+ data accuracy across 10+ AI initiatives and built internal tools that save 150+ hours daily for a 200-person team. Seeking ML Engineer, GenAI Engineer, or Data Scientist roles.

TECHNICAL SKILLS

ML/DL: PyTorch, TensorFlow, Scikit-Learn, XGBoost, CNNs, Transfer Learning, Feature Engineering, Hyperparameter Tuning

GenAI: LangChain, LangGraph, RAG Pipelines, LLM Fine-tuning (LoRA/QLoRA), Hugging Face, Agentic AI, RLHF, Guardrails, Vector DBs (FAISS/ChromaDB)

AWS/MLOps: SageMaker, Bedrock, Lambda, Docker, Kubernetes, MLflow, CI/CD, Model Monitoring, TorchServe

Python & SQL: NumPy, Pandas, FastAPI, Pytest, Streamlit, Plotly | PostgreSQL, Window Functions, CTEs, Airflow, Spark

Tools: Git, Linux, Docker, Kubernetes, Tableau, Power BI, Weights & Biases, Agile/Scrum

WORK EXPERIENCE

AMAZON, Hyderabad

March 2026 – Present

AI Benchmarking Specialist

- Building and deploying an agentic AI assistant for Amazon sellers — a conversational system that queries Seller Central data (inventory, orders, revenue, shipping) and delivers business insights in natural language with inline visualizations
- Designed the full RAG pipeline: SP-API data ingestion → ETL → PostgreSQL (structured) + ChromaDB (vector embeddings) → LangGraph agent with SQL and retrieval tools → Claude/GPT response generation
- Implemented safety and guardrails layer — query classification, harmful content detection, policy-bypass prevention, and hallucination reduction achieving 99% response accuracy
- Leading RLHF evaluation workflows: rating helpfulness, factual accuracy, and data correctness across 1000+ query-response pairs to continuously improve model quality
- Benchmarking AI models end-to-end — defining test plans, evaluation metrics, and acceptance criteria across accuracy, robustness, bias, and reliability

AMAZON WEB SERVICES (AWS), Bengaluru

October 2024 – March 2026

Machine Learning Data Associate II

- Built automated data pipelines using Python (Pandas, NumPy) and SQL to prepare, validate, and annotate large-scale datasets across text, code (QAL), and computer vision (image/video) — maintaining 99%+ accuracy across all deliverables
- Engineered high-precision training datasets for Agentic AI models across Finance, Mortgage, Health, and Legal domains — 10+ initiatives including Red-Teaming, mathematical evaluation, and multi-turn conversation flows
- Developed “SageMaker Smart Stats” from scratch — a JavaScript-based productivity analytics tool that automated metric tracking, surfaced workflow bottlenecks, and saved ~45 minutes per associate daily across a 200-person team (15–20% output increase)
- Authored technical SOPs and annotation guidelines for complex ML tasks — reduced guideline-related errors by 15% and significantly cut rework cycles, becoming the go-to documentation resource for the team
- Built custom Tableau and Python dashboards for data quality monitoring — flagging annotation drift, inter-annotator agreement scores, and SLA adherence in real-time

AMAZON, Delhi (Remote)

September 2023 – March 2024

Artificial Intelligence Associate

- Built Python-based automation scripts for data validation and quality scoring — improved AI model training accuracy from 94% to 99% while reducing manual rework by 30%
- Developed performance monitoring dashboards using Python (Pandas, Matplotlib) and Excel VBA — tracking per-associate throughput, error rates, and SLA adherence in real-time
- Maintained 99.5% dataset accuracy across all deliverables with zero SLA misses — identified systematic labeling errors through statistical analysis and implemented corrective feedback loops

Educator — Physics, Mathematics, Python, SQL

- Taught Physics, Mathematics, Python programming, and SQL fundamentals to 500+ students over 7 years — average student scores above 91%
- Introduced Python and SQL modules into the curriculum — covering data types, loops, functions, Pandas basics, and writing queries for real-world datasets
- Taught probability, statistics, and linear algebra — the mathematical foundations that directly map to machine learning and data science
- Tracked student performance using spreadsheets and basic Python scripts — identified weak areas early and adjusted teaching approach accordingly
- Managed batches of 80+ students with structured lesson plans, regular assessments, and personalized feedback cycles

PROJECTS

GenAI Multimodal Chatbot — AWS Bedrock

Full-stack generative AI chatbot handling text, images, and documents. Built on Claude via AWS Bedrock with streaming responses, conversation memory, and enterprise guardrails. Serves 500+ concurrent users.

Tech: AWS Bedrock, Claude, Python, Streamlit, DynamoDB

RAG Document QA Pipeline — AWS Bedrock

Production-grade RAG pipeline with semantic chunking, FAISS vector store, and Bedrock embeddings. Answers queries with proper source citations. Handles enterprise document loads.

Tech: LangChain, AWS Bedrock, FAISS, FastAPI, Python

End-to-End ML Pipeline — SageMaker (Fraud Detection)

Complete MLOps pipeline: automated feature engineering, hyperparameter tuning, A/B testing, and real-time inference endpoint. Achieved 96.3% precision on production data.

Tech: SageMaker, XGBoost, Feature Store, Lambda, Step Functions

LLM Fine-tuning with PyTorch (LoRA/QLoRA)

Parameter-efficient fine-tuning of Llama-2 for domain-specific financial QA. 4-bit quantization for consumer GPU deployment. 23% accuracy improvement over base model.

Tech: PyTorch, LoRA, Hugging Face, PEFT, bitsandbytes

Sentiment Analysis API — Transformers

Real-time sentiment classification using fine-tuned DistilBERT. Containerized with Docker, deployed on AWS Lambda. Sub-100ms inference latency, handles 1000+ requests/min.

Tech: Transformers, FastAPI, Docker, Lambda, DistilBERT

Real-Time Stock Anomaly Detector

Pulls live market data via API, applies Z-score and rolling-window anomaly detection. Clean OOP structure with type hints and full test coverage. Interactive Plotly visualizations.

Tech: Python, Pandas, NumPy, Plotly, Statistics, OOP

E-Commerce SQL Analytics Engine

Advanced analytics on real e-commerce data: cohort retention, RFM customer segmentation, funnel conversion tracking, revenue forecasting with window functions. EXPLAIN plan optimization.

Tech: PostgreSQL, Window Functions, CTEs, Query Optimization, RFM

CERTIFICATIONS

AWS Certified AI Practitioner (Dec 2025) | **AWS Generative AI Developer** (2026) | **Google Data Analytics Professional** (Apr 2025)

Python for Data Science — IBM (2024) | **SQL for Data Science — UC Davis** (2024) | **Project Management: Agile/Scrum** (Jan 2024)

EDUCATION & AWARDS

B.Tech in Electrical and Electronics Engineering — National Institute of Technology Sikkim, Ravangla 2012 – 2016

Prime Player Award (Feb & Apr 2025) — Highest productivity/reliability across entire team at AWS. Twice in 3 months.

Leadership Recognition — Shout-outs for Agentic AI SOPs, multilingual datasets, exceeding expectations.